
Problem set #4

Release 0.1

Brian Granger, John Hunter and Fernando Pérez.

February 10, 2010

Contents

| | | |
|----------|----------------------------|-----------|
| 1 | FFT Image Denoising | i |
| 2 | Hints | ii |

1 FFT Image Denoising

Illustrates: 2-d image denoising, use of the numpy FFT library, array manipulations, image plotting.

The convolution of an input with with a linear filter in the temporal or spatial domain is equivalent to multiplication by the Fourier transforms of the input and the filter in the spectral domain. This provides a conceptually simple way to think about filtering: transform your signal into the frequency domain, dampen the frequencies you are not interested in by multiplying the frequency spectrum by the desired weights, and then apply the inverse transform to the modified spectrum, back into the original domain. In the example below, we will simply set the weights of the frequencies we are uninterested in (the high frequency noise) to zero rather than dampening them with a smoothly varying function. Although this is not usually the best thing to do, since sharp edges in one domain usually introduce artifacts in another (eg high frequency “ringing”), it is easy to do and sometimes provides satisfactory results.

After reading the image file `moonlanding.png`, try to produce images like the ones in the figure. We will describe the process here and provide some hints as to what you need to think about.

The image in the upper left panel of the Figure is a grayscale photo of the moon landing. There is a banded pattern of high frequency noise polluting the image. In the upper right panel we see the 2D spatial frequency spectrum. The FFT output in the `numpy.fft` module is packed with the lower frequencies starting in the upper left, and proceeding to higher frequencies as one moves to the center of the spectrum (this is the most efficient way numerically to fill the output of the

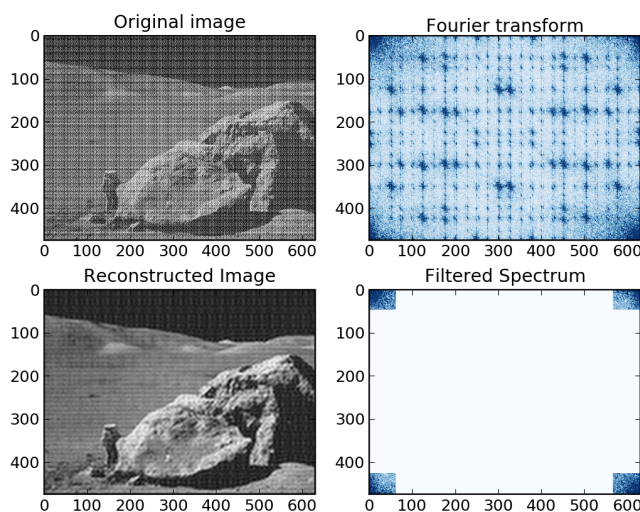


Figure 1: High frequency noise filtering of a 2D image in the Fourier domain. The upper panels show the original image (left) and spectral power (right) and the lower panels show the same data with the high frequency power set to zero. Although the input and output images are grayscale, you can provide colormaps to `imshow()` to plot them in pseudo-color, with the `cmap` argument to `imshow()`, which accepts any of the colormaps found in the `matplotlib.cm` module.

FFT algorithm). Because the input signal is real, the output spectrum is complex and symmetrical: the transformation values beyond the midpoint of the frequency spectrum (the Nyquist frequency) correspond to the values for negative frequencies and are simply the mirror image of the positive frequencies below the Nyquist (this is true for the 1D, 2D and ND FFTs in `numpy`).

You should compute the 2D spatial frequency spectra of the luminance image, zero out the high frequency components, and inverse transform back into the spatial domain. You can plot the input and output images with `plt.imshow()`, but you should observe that if you show the power spectrum (the absolute value of the FFT) directly, you will only see white, and not the image in the Figure's upper right panel. This is due to the fact that the power spectrum has a small number of pixels with extremely high amplitude, which completely swamp the contrast (you can verify this by playing with a histogram of the data). You will thus need to clip the limits of the color range to a small range of values where most of the power actually lives (say 95%).

2 Hints

- The upper right panel is obtained after finding the range of values that contain 95% of the power. For this, the `mlab.prctile()` function will be useful.
- The `numpy.fft` module contains the necessary FFT routines for this exercise.
- In Python, a complex number `z` has `z.real` and `z.imag` attributes for its real and imaginary parts.